

Analyze_AB_Test_Results

November 11, 2020

1 Analyze A/B Test Results

Table of Contents 1. Introduction 2. Part I - Probability 3. Part II - A/B Test 4. Part III - Regression

Introduction A/B tests are very commonly performed by data analysts and data scientists. It is important that we get some practice working with the difficulties of these tests.

For this project, we will be working to understand the results of an A/B test run by an e-commerce website. our goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

```
[1]: # importing the needed libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import random
import timeit
%matplotlib inline
```

Part I - Probability

```
[2]: random.seed(42)
```

```
[3]: # reading the "ab_data" dataset
AB_data= pd.read_csv('ab_data.csv')
```

```
[4]: AB_data.head(10)
```

```
[4]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1
5	936923	2017-01-10 15:20:49.083499	control	old_page	0
6	679687	2017-01-19 03:26:46.940749	treatment	new_page	1

7	719014	2017-01-17 01:48:29.539573	control	old_page	0
8	817355	2017-01-04 17:58:08.979471	treatment	new_page	1
9	839785	2017-01-15 18:11:06.610965	treatment	new_page	1

user_id: the id associated with each user and it's unique; hence no two users share the same id.
timestamp: is the date and time as when the user accessed the page. **group**: 1- **treatment**, which are the people who got experimented on using the new page. 2- **control**, which are the people who kept using the old page, and they're used as a baseline measure. **landing_page**: the old and the new page. **converted**: 0 for the users who did not pay for the company's product, and 1 for who did.

```
[5]: #the number of rows
all_users=AB_data.shape[0]
all_users
```

```
[5]: 294478
```

```
[6]: AB_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   user_id         294478 non-null  int64
1   timestamp       294478 non-null  object
2   group           294478 non-null  object
3   landing_page    294478 non-null  object
4   converted       294478 non-null  int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB
```

```
[7]: # the number of unique users in the dataset
unique_users= AB_data['user_id'].nunique()
unique_users
```

```
[7]: 290584
```

```
[8]: # getting the number of users converted
converted_users= (AB_data.query('converted==1')['user_id']).nunique()
converted_users
```

```
[8]: 35173
```

```
[9]: # the proportion of users converted
prop_converted_users= converted_users/unique_users
prop_converted_users
```

```
[9]: 0.12104245244060237
```

```
[10]: print("The proportion of the converted users equals approximately {0: .0%}".  
      ↪format(prop_converted_users))
```

The proportion of the converted users equals approximately 12%

```
[11]: # The number of times the new_page and treatment don't line up  
NP_trmnt_NoLineUp=AB_data.query('(landing_page != "new_page" and group ==  
    ↪"treatment") or (landing_page == "new_page" and group !=  
    ↪"treatment")')['user_id'].count()  
NP_trmnt_NoLineUp
```

```
[11]: 3893
```

```
[12]: # missing values  
if (AB_data.isnull().values.any()) == False:  
    print("There's no missing value in ab_data dataset.")  
elif (AB_data.isnull().values.any()) == True:  
    null_num= AB_data.isnull().sum()  
    if (null_num==1):  
        print("There's {} missing value in ab_data dataset.".format(null_num))  
    elif (null_num>1):  
        print("There're {} missing values in ab_data dataset.".format(null_num))
```

There's no missing value in ab_data dataset.

```
[13]: AB_data_new= AB_data.copy()
```

I made a copy to keep my original data intact, in case anything went wrong during the analysis process

```
[14]: AB_data_new.drop(AB_data_new.query('(landing_page != "new_page" and group ==  
    ↪"treatment") or (landing_page == "new_page" and group != "treatment") or  
    ↪(landing_page != "old_page" and group == "control") or (landing_page ==  
    ↪"old_page" and group != "control")').index, inplace=True)
```

I deleted the rows where treatment is not aligned with new_page or control is not aligned with old_page, because we cannot be sure if these rows truly received the new or old page. Hence it's useless to me.

```
[15]: # making sure that the dropping process went well by subtracting the number of  
    ↪all users from our new updated dataset; the result should equal  
    ↪NP_trmnt_NoLineUp which I already calculated  
check= all_users - AB_data_new.shape[0]  
check
```

[15]: 3893

```
[16]: # making sure that the dropped rows were the intended ones, the result of the
      ↪code should be 0
      AB_data_new[((AB_data_new.landing_page == 'new_page')== (AB_data_new.group ==
      ↪'treatment')) == False].shape[0]
```

[16]: 0

The dropping process went very well. The rows that has mismatching were eliminated successfully.

```
[17]: # the number of unique user_ids in AB_data_new
      AB_data_new['user_id'].nunique()
```

[17]: 290584

```
[18]: AB_data_new.shape[0]
```

[18]: 290585

As it shows from the previous two cells that there's one repeated user id.

```
[19]: #the duplicated user_id with the row information
      AB_data_new[AB_data_new.duplicated(subset=['user_id'], keep= False)]
```

```
[19]:      user_id      timestamp      group landing_page  converted
1899   773192  2017-01-09 05:37:58.781806  treatment      new_page         0
2893   773192  2017-01-14 02:55:59.590927  treatment      new_page         0
```

```
[20]: # getting rid of one of the duplicates
      AB_data_new.drop(AB_data_new[(AB_data_new['user_id'] == 773192) &
      ↪(AB_data_new['timestamp'] == '2017-01-09 05:37:58.781806')].index,
      ↪inplace=True)
```

I got rid of one of the duplicates since the only difference between them is the timestamp, and this difference does not affect my work.

```
[21]: # checking the success of the dropping process
      AB_data_new[(AB_data_new['user_id']==773192)]
```

```
[21]:      user_id      timestamp      group landing_page  converted
2893   773192  2017-01-14 02:55:59.590927  treatment      new_page         0
```

As we can see, there's only one user with the id **773192**.

```
[22]: # the number of converted users in AB_data_new
      new_converted_users= AB_data_new.query('converted == 1')['user_id'].nunique()
      new_converted_users
```

[22]: 34753

```
[23]: #the probability of an individual converting regardless of the page they receive  
new_prob_converted_users= new_converted_users/AB_data_new.shape[0]
```

```
[24]: print("The probability of an individual converting regardless of the page they_  
↪receive equals {:.2%}".format(new_prob_converted_users))
```

The probability of an individual converting regardless of the page they receive equals 11.96%

```
[25]: #the users who converted while they're in the control group  
control_converted_users= AB_data_new.query('group == "control" and converted ==_  
↪1')['user_id'].nunique()  
control_converted_users
```

[25]: 17489

```
[26]: # all the users in the control group  
control_users= AB_data_new.query('group == "control"')['user_id'].nunique()  
control_users
```

[26]: 145274

```
[27]: #the probability that users converted while they were in the control group  
prob_control_converted_users= control_converted_users/control_users
```

```
[28]: print("The probability that users converted while they were in the control_  
↪group equals {:.2%}".format(prob_control_converted_users))
```

The probability that users converted while they were in the control group equals 12.04%

```
[29]: #the users who converted while they're in the treatment group  
treatment_converted_users= AB_data_new.query('group == "treatment" and_  
↪converted == 1')['user_id'].nunique()  
treatment_converted_users
```

[29]: 17264

```
[30]: # all the users in the treatment group  
treatment_users= AB_data_new.query('group == "treatment"')['user_id'].nunique()  
treatment_users
```

[30]: 145310

```
[31]: #the probability that users converted while they were in the treatment group
prob_treatment_converted_users= treatment_converted_users/treatment_users
```

```
[32]: print("The probability that users converted while they were in the treatment_
      ↪group equals {:.2%}".format(prob_treatment_converted_users))
```

The probability that users converted while they were in the treatment group equals 11.88%

```
[33]: # the numbers of users who received the new page
NewPage_users= AB_data_new.query('landing_page == "new_page"')['user_id'].
      ↪nunique()
NewPage_users
```

```
[33]: 145310
```

```
[34]: # the probability that an individual received the new page
prob_NewPage= NewPage_users/AB_data_new.shape[0]
prob_NewPage
```

```
[34]: 0.5000619442226688
```

```
[35]: print("The probability that an individual received the new page equals {:.2%}".
      ↪format(prob_NewPage))
```

The probability that an individual received the new page equals 50.01%

Notwithstanding the page that the users received, the probability of them converting is 11.96%. And if we took into consideration that the users were in the control group (hence received the old page), the probability of them converting is 12.04%. Moreover, if we took into account that the users were in the treatment group (hence received the new page), the probability of them converting is 11.88%. If we look at this point at the numbers, we would probably say that the old page did a better job making the users convert, though it's a negligible convert. **Yet**, we can hardly consider these numbers a sufficient evidence; because the results are pretty close to each other and couldn't help us decide whether any page would lead to more conversions or not.

Part II - A/B Test

For now, let's consider for the sake of argument that I need to make the decision just based on all the data provided. If I want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, the null hypotheses would be **H0**: $p_{new} - p_{old} \leq 0$, and the alternative hypotheses would be **H1**: $p_{new} - p_{old} > 0$.

Let's assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, let's assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page. we're going to perform the sampling distribution for the difference in converted between the two pages over 10,000 iterations of calculating an estimate from the null:

```
[36]: #the convert rate for under the null
p_new= AB_data_new.query('converted == 1')['user_id'].nunique()/AB_data_new.
      ↪user_id.nunique()
p_new
```

```
[36]: 0.11959708724499628
```

```
[37]: #the convert rate for under the null
p_old= AB_data_new.query('converted == 1')['user_id'].nunique()/AB_data_new.
      ↪user_id.nunique()
p_old
```

```
[37]: 0.11959708724499628
```

```
[38]: #the number of unique users who has new page using AB_data_new dataframe
n_new= AB_data_new.query('landing_page == "new_page"')['user_id'].nunique()
n_new
```

```
[38]: 145310
```

```
[39]: #the number of unique users who has old page using AB_data_new dataframe
n_old= AB_data_new.query('landing_page == "old_page"')['user_id'].nunique()
n_old
```

```
[39]: 145274
```

```
[40]: #Simulate n transactions with a convert rate of p under the null
new_page_converted= np.random.choice([0,1], n_new, p=(p_new, 1-p_new))
new_page_converted
```

```
[40]: array([1, 1, 1, ..., 1, 1, 1])
```

```
[41]: #Simulate n transactions with a convert rate of p under the null
old_page_converted= np.random.choice([0,1], n_old, p=(p_old, 1-p_old))
old_page_converted
```

```
[41]: array([1, 1, 0, ..., 1, 1, 1])
```


```
[42]: # - for the simulated values in the previous two cells
      (new_page_converted.mean())-(old_page_converted.mean())
```

```
[42]: -0.0008170246352350308
```

```
[43]: #Simulate 10,000 - values like what we just did
new_converted_simulation = np.random.binomial(n_new, p_new, 10000)/n_new
old_converted_simulation = np.random.binomial(n_old, p_old, 10000)/n_old
p_diffs = new_converted_simulation - old_converted_simulation
```

```
[44]: #Plotting a histogram of the p_diffs
plt.hist(p_diffs, color='crimson')
```

```
[44]: (array([ 23., 181., 744., 1931., 2881., 2482., 1272., 404., 76.,
        6.]),
array([-0.00429, -0.00339045, -0.0024909, -0.00159135, -0.0006918,
        0.00020775, 0.0011073, 0.00200685, 0.0029064, 0.00380595,
        0.0047055 ]),
<BarContainer object of 10 artists>)
```



The proportion of the p_diffs that are greater than the actual difference observed in ab_data.csv:

```
[45]: # distribution under the null hypothesis
null_v= np.random.normal(0, p_diffs.std(), p_diffs.size)
```

```
[46]: # number of users who converted in the new page and the old page
new_page_converted_users= AB_data_new.query('converted == 1 and landing_page == "new_page"')['user_id'].nunique()
old_page_converted_users= AB_data_new.query('converted == 1 and landing_page == "old_page"')['user_id'].nunique()
```

```
[47]: # actual convert rate
new_actual_convert= new_page_converted_users/n_new
old_actual_convert= old_page_converted_users/n_old
```




```
[48]: #observed difference in converted rate
convert_observed_difference= new_actual_convert - old_actual_convert
convert_observed_difference
```

```
[48]: -0.0015782389853555567
```

```
[49]: # the proportion of the p_diffs that are greater than the actual difference
      ↳observed in ab_data.csv
plt.hist(p_diffs, color='crimson')
plt.axvline(x=convert_observed_difference,color='greenyellow')
```

```
[49]: <matplotlib.lines.Line2D at 0x7f8afec25460>
```



```
[50]: # the proportion of the p_diffs that are greater than the actual difference
      ↳observed in ab_data.csv
p_value1=(null_v >convert_observed_difference).mean()
p_value1
```

```
[50]: 0.9076
```

In scientific studies, there's a value called the p-value. If we assumed that the null hypothesis was true and the p-value was less than 0.05, that would indicate there's a low probability in finding values remarkably greater-than, less-than or equal to `convert_observed_difference`. Since `p_old` (or `p_new`) does not fall within the critical region, we fail to reject the null hypothesis (`p_old > Alpha` → `0.11959708724499628 > 0.05`). Therefore, with a Type I error rate of 5% there's a higher probability of conversion in the old page than the new page. Also, the `p_value` is greater than

Alpha, hence we accept the null hypothesis. P-values are the probability of obtaining an effect at least as extreme as the one in our sample data, assuming the truth of the null hypothesis.

source: <https://blog.minitab.com/blog/adventures-in-statistics-2/understanding-hypothesis-tests-significance-levels-alpha-and-p-values-in-statistics>

```
[51]: # using a built-in to achieve similar results to what we did above.
import statsmodels.api as sm

convert_old = old_page_converted_users
convert_new = new_page_converted_users
n_old=n_old
n_new=n_new
```

```
[52]: # using stats.proportions_ztest to compute my test statistic and p-value
Ztest_score, p_value2 = sm.stats.proportions_ztest(np.
    ↳ array([convert_new,convert_old]),np.array([n_new,n_old]), alternative = '
    ↳ 'larger')
```

```
[53]: Ztest_score, p_value2
```

```
[53]: (-1.3109241984234394, 0.9050583127590245)
```

The minus sign before the `Ztest_score` value indicates that the value is below the mean. Since the `Ztest_score` (which is 1.3109241984234394) is not greater than 1.96 (which is the critical value with significance of 5%), we fail to reject the null hypothesis; which means that the old page has a better converted rate than the new page or equal to it. Moreover, the `p-value` (which is 0.9050583127590245) is not less than 0.05, and it is nearly similar to the p-value that we calculated previously (which is 0.9026). And these conclusions prove that the z-test is consistent with what we achieved so far.

Source: https://www.google.com/url?sa=t&rct=j&q=&esrc=s&source=web&cd=&ved=2ahUKEwj7f_rkvPsAh

Part III - A regression approach

Since each row is either a conversion or no conversion, Logistic Regression would be the ideal type of regression in this case. And also because we are interested in knowing the odds of conversion, and Logistic Regression works with the categorical outcomes.

```
[54]: #creating a column for the intercept
AB_data_new['intercept'] = 1
```

we need an intercept term so that our regression output does not automatically predict probability 1/2 when our variables are all zero. source: <https://www.quora.com/Why-do-we-add-a-column-of-1s-intercept-in-multiple-logistic-regression>

```
[55]: #creating a dummy variable column for which page each user received
AB_data_new= AB_data_new.join(pd.get_dummies(AB_data_new['landing_page']))
AB_data_new.head()
```

```
[55]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	intercept	new_page	old_page
0	1	0	1
1	1	0	1
2	1	1	0
3	1	1	0
4	1	0	1

```
[56]: # adding an ab_page column, which is 1 when an individual receives the
      ↪ treatment and 0 if control
AB_data_new['ab_page'] = pd.get_dummies(AB_data['group']) ['treatment']
AB_data_new.head()
```

```
[56]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	intercept	new_page	old_page	ab_page
0	1	0	1	0
1	1	0	1	0
2	1	1	0	1
3	1	1	0	1
4	1	0	1	0

I'm going to use statsmodels to import my regression model. Instantiate the model, and fit the model using the two columns I created above to predict whether or not an individual converts:

```
[57]: # Creating Logistic regression model for ab_page and the converted variable
Logistic_model1 = sm.Logit(AB_data_new['converted'],
      ↪AB_data_new[['intercept', 'ab_page']])
```

```
[58]: Results1 = Logistic_model1.fit()
```

```
Optimization terminated successfully.
      Current function value: 0.366118
      Iterations 6
```

```
[59]: print(Results1.summary())
```

Logit Regression Results

=====						
Dep. Variable:	converted		No. Observations:	290584		
Model:	Logit		Df Residuals:	290582		
Method:	MLE		Df Model:	1		
Date:	Wed, 11 Nov 2020		Pseudo R-squ.:	8.077e-06		
Time:	21:51:19		Log-Likelihood:	-1.0639e+05		
converged:	True		LL-Null:	-1.0639e+05		
Covariance Type:	nonrobust		LLR p-value:	0.1899		
=====						
	coef	std err	z	P> z	[0.025	0.975]

intercept	-1.9888	0.008	-246.669	0.000	-2.005	-1.973
ab_page	-0.0150	0.011	-1.311	0.190	-0.037	0.007
=====						

As we can see, The p-value associated with `ab_page` is 0.190. This Logit Regression model is attempting to predict whether a user will convert based on whether their page is new or old. The null hypotheses is that there's no difference between the two groups (treatment and control). While the Alternative hypotheses is that there's difference. In Part II the assumption is that the old page is better except if the new page proves to be better at a **Type I** error rate of 5%. This part has different explanatory variable or factor for the result in comparison to Part II.

Meanwhile, let's consider other things that might influence whether or not an individual converts; such as the time (timestamp variable). We can look and see if the converted rate is determined by a particular time or day when the users are using the website. And we can do so by turning the `timestamp` variable into a categorical variable, where it segments the time into **morning**, **afternoon**, **evening** and **past midnight**, or into **weekends** and **weekdays**. As helpful and productive as this might seem, by adding additional terms into our regression model we're adding up to the complexity of interpretation. Also, the added terms might not affect the outcome, or it might affect the other participating variables; which might result in erroneous outcomes.

Now along with testing if the conversion rate changes for different pages, we'll also add an effect based on which country a user lives. But first, we'll need to read in the `countries.csv` dataset and merge together our datasets on the appropriate rows:

```
[60]: # reading the "countries" dataset
countries= pd.read_csv('countries.csv')
countries.head(10)
```

```
[60]:  user_id country
0    834778      UK
1    928468      US
2    822059      UK
3    711597      UK
4    710616      UK
5    909908      UK
6    811617      US
7    938122      US
```

```
8 887018 US
9 820683 US
```

```
[61]: # merging the two datasets on user_id
ABData_and_countries= AB_data_new.merge(countries, on ='user_id', how='left')
ABData_and_countries.head(10)
```

```
[61]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	
5	936923	2017-01-10 15:20:49.083499	control	old_page	0	
6	679687	2017-01-19 03:26:46.940749	treatment	new_page	1	
7	719014	2017-01-17 01:48:29.539573	control	old_page	0	
8	817355	2017-01-04 17:58:08.979471	treatment	new_page	1	
9	839785	2017-01-15 18:11:06.610965	treatment	new_page	1	

	intercept	new_page	old_page	ab_page	country
0	1	0	1	0	US
1	1	0	1	0	US
2	1	1	0	1	US
3	1	1	0	1	US
4	1	0	1	0	US
5	1	0	1	0	US
6	1	1	0	1	CA
7	1	0	1	0	US
8	1	1	0	1	UK
9	1	1	0	1	CA

```
[62]: # getting the countries in country column to create the dummy variables
countries.country.unique()
```

```
[62]: array(['UK', 'US', 'CA'], dtype=object)
```

```
[63]: # creating the dummy variables
ABData_and_countries[['CA','UK','US']] = pd.
↳get_dummies(ABData_and_countries['country'])
```

```
[64]: ABData_and_countries.head(10)
```

```
[64]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	

4	864975	2017-01-21	01:52:26.210827	control	old_page	1
5	936923	2017-01-10	15:20:49.083499	control	old_page	0
6	679687	2017-01-19	03:26:46.940749	treatment	new_page	1
7	719014	2017-01-17	01:48:29.539573	control	old_page	0
8	817355	2017-01-04	17:58:08.979471	treatment	new_page	1
9	839785	2017-01-15	18:11:06.610965	treatment	new_page	1

	intercept	new_page	old_page	ab_page	country	CA	UK	US
0	1	0	1	0	US	0	0	1
1	1	0	1	0	US	0	0	1
2	1	1	0	1	US	0	0	1
3	1	1	0	1	US	0	0	1
4	1	0	1	0	US	0	0	1
5	1	0	1	0	US	0	0	1
6	1	1	0	1	CA	1	0	0
7	1	0	1	0	US	0	0	1
8	1	1	0	1	UK	0	1	0
9	1	1	0	1	CA	1	0	0

```
[65]: Logistic_model2 = sm.Logit(ABData_and_countries['converted'],
    ↪ABData_and_countries[['intercept', 'new_page', 'UK', 'US']])
```

```
[66]: Results2 = Logistic_model2.fit()
```

```
Optimization terminated successfully.
    Current function value: 0.366113
    Iterations 6
```

```
[67]: print(Results2.summary())
```

```

                        Logit Regression Results
=====
Dep. Variable:          converted    No. Observations:          290584
Model:                  Logit       Df Residuals:              290580
Method:                  MLE        Df Model:                  3
Date:                   Wed, 11 Nov 2020    Pseudo R-squ.:          2.323e-05
Time:                   21:51:20    Log-Likelihood:         -1.0639e+05
converged:              True        LL-Null:                 -1.0639e+05
Covariance Type:        nonrobust    LLR p-value:            0.1760
=====

```

	coef	std err	z	P> z	[0.025	0.975]
intercept	-2.0300	0.027	-76.249	0.000	-2.082	-1.978
new_page	-0.0149	0.011	-1.307	0.191	-0.037	0.007
UK	0.0506	0.028	1.784	0.074	-0.005	0.106
US	0.0408	0.027	1.516	0.130	-0.012	0.093

```
=====
```

Again, based on the p-value shown it does not appear that the country had an impact on the conversion.

Though we have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there's significant effects on conversion.

```
[68]: #Creating a new interacton variable between new_page and country US, UK and CA
ABData_and_countries['UK_new_page'] = ABData_and_countries['new_page'] *
↳ABData_and_countries['UK']
ABData_and_countries['US_new_page'] = ABData_and_countries['new_page'] *
↳ABData_and_countries['US']
ABData_and_countries['CA_new_page'] = ABData_and_countries['new_page'] *
↳ABData_and_countries['CA']
```

```
[69]: ABData_and_countries.head()
```

```
[69]:  user_id      timestamp      group landing_page  converted  \
0   851104  2017-01-21 22:11:48.556739   control    old_page         0
1   804228  2017-01-12 08:01:45.159739   control    old_page         0
2   661590  2017-01-11 16:55:06.154213  treatment    new_page         0
3   853541  2017-01-08 18:28:03.143765  treatment    new_page         0
4   864975  2017-01-21 01:52:26.210827   control    old_page         1

   intercept  new_page  old_page  ab_page  country  CA  UK  US  UK_new_page  \
0           1         0         1         0       US   0   0   1             0
1           1         0         1         0       US   0   0   1             0
2           1         1         0         1       US   0   0   1             0
3           1         1         0         1       US   0   0   1             0
4           1         0         1         0       US   0   0   1             0

   US_new_page  CA_new_page
0             0             0
1             0             0
2             1             0
3             1             0
4             0             0
```

```
[70]: #Create logistic regression
Logistic_model3 = sm.Logit(ABData_and_countries['converted'],
↳ABData_and_countries[['intercept', 'new_page', 'US_new_page',
↳'UK_new_page', 'UK', 'US']])
```

```
[71]: Results3 = Logistic_model3.fit()
```

Optimization terminated successfully.

Current function value: 0.366109

Iterations 6

```
[72]: print(Results3.summary())
```

```

                        Logit Regression Results
=====
Dep. Variable:          converted    No. Observations:          290584
Model:                  Logit      Df Residuals:              290578
Method:                  MLE       Df Model:                  5
Date:                   Wed, 11 Nov 2020    Pseudo R-squ.:          3.482e-05
Time:                   21:51:21    Log-Likelihood:         -1.0639e+05
converged:               True        LL-Null:               -1.0639e+05
Covariance Type:        nonrobust    LLR p-value:            0.1920
=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
intercept    -2.0040     0.036   -55.008     0.000     -2.075     -1.933
new_page     -0.0674     0.052   -1.297     0.195     -0.169      0.034
US_new_page   0.0469     0.054    0.872     0.383     -0.059      0.152
UK_new_page   0.0783     0.057    1.378     0.168     -0.033      0.190
UK            0.0118     0.040    0.296     0.767     -0.066      0.090
US            0.0175     0.038    0.465     0.642     -0.056      0.091
=====
```

```
[73]: np.exp(Results3.params)
```

```
[73]: intercept    0.134794
      new_page     0.934776
      US_new_page  1.048001
      UK_new_page  1.081428
      UK           1.011854
      US           1.017682
      dtype: float64
```

The new_page's p-value is 0.195 which is quite similar to the previous ones. As we can see from the **Logit Regression Results**, the new_page's coefficient differs from US_new_page's and UK_new_page's. Also, z-score for the variables weren't large enough for predicting conversion rate, and the p-values are greater than 0.05 except for the intercept. For that reason, the location of the user who's using the page is irrelevant in terms of conversion rate. Moreover, the probability of conversion is 7.0% for every unit new_page decreases, holding all other variables constant. Also, the likelihood of conversion for UK and new_page's users is 1.08 higher than CA and new_page's users, and 1.18% for the users in UK than CA, holding all other variables constant. Additionally, the likelihood of conversion for US and new_page's users is 1.04 higher than CA and new_page's users, and 1.75% for the users in US than CA, holding all other variables constant. In the end, we don't have enough evidence to reject the null hypothesis based on our A/B testing. So under these circumstances, there's no point in switching to the new page, when the old page is working just as well.